

ω Cen: kinematic data preparation for the dark-matter inference

oCen_dm project notes

2026 September 21

Abstract

This note documents how the proper-motion and line-of-sight dispersion profiles that enter the ω Cen mass inference were assembled: which catalogues, which quality selections, which estimator, and which cross-checks. The central methodological point is that both major proper-motion catalogues contain a large population of stars whose quoted uncertainties are underestimated, and that any dispersion measured from them is a statement about the error model rather than about the cluster. Once each survey’s own quality selection is respected, HST and *Gaia* EDR3 agree in their 300–360'' overlap at $\sigma_{\text{Gaia}}/\sigma_{\text{HST}} = 1.03 \pm 0.13$, the precision being set by the 17 quality-flagged *Gaia* stars available there; a Pristine-based measurement using a different membership definition agrees with *Gaia* to 0.5 ± 1.2 per cent, though it is not an independent astrometric check because its proper motions are the same *Gaia* measurements. Both published profiles are reproduced to a few tenths of a per cent. The current dataset is 89 binned points from five profiles spanning approximately 0.1–55 pc at the rung-0 distance. The current selections are tested, but correlated systematics, tracer differences and the outer equilibrium assumption remain open.

1 Scope and principle

The inference asks whether ω Cen requires a dark component beyond its stars and remnants. That question is decided in the outskirts, where a halo would raise the dispersion above what the luminous mass supports. The outskirts are also where the data are worst. Every methodological choice below follows from a single principle, which we arrived at independently and later found stated explicitly by Häberle et al. (2025):

A star whose measurement uncertainty approaches the velocity dispersion being measured makes the inferred dispersion a function of the assumed error model rather than of the cluster.

Such stars are therefore excluded rather than modelled.

2 Data sources

Dataset	Source	Content	Role
oMEGACat VI	Häberle et al. (2025)	1.48×10^6 HST PMs to 466''	primary inner PM
oMEGACat I / MUSE	Häberle et al. (2025)	3×10^5 LOS velocities	inner LOS
<i>Gaia</i> EDR3 members	Vasiliev & Baumgardt (2021)	2.28×10^5 stars to 0.67°	primary outer PM
<i>Gaia</i> DR2 profile	Baumgardt et al. (2019)	9 binned points	outer PM, published
Pristine periphery	Kuzma & Ishigaki (2025)	1.57×10^5 stars to 5.1°	external check
Periphery spectroscopy	Kuzma et al. (2026)	592 LOS velocities to 3.15°	external check

Table 1: Catalogues used. Only the first three enter the current likelihood; *Gaia* DR2 remains a legacy comparison.

3 The cluster centre

Every radius in this analysis depends on one pair of numbers, and until 2026-09-19 they were wrong. Four modules hard-coded (201.696833, −47.476583) with the comment “Baumgardt catalogue centre”; the directory that comment cites was never downloaded, and the value was written from memory. It sits 10.7'' north of the centre oMEGACat itself uses and 9.9'' north of Vasiliev & Baumgardt’s. Those two agree with each other to 0.83''. The symptom: stars labelled as lying 2.6–3.2'' from the centre were really at 7.8–13.9''.

The centre is now derived rather than written down. oMEGACat’s astrometry is tied to a pixel grid whose reference pixel carries the Anderson & van der Marel photometric centre; fitting the local solution from the 272 stars within 60 px gives (201.696833, −47.479569), and the recovered origin is stable against the fitting window and against adding quadratic terms. Vasiliev’s member table inverts to (201.696838, −47.479339). We adopt the HST value. Six tests pin it, one of which greps the source tree so the literal cannot return.

Where it mattered: **nowhere for *Gaia***, every bin within 0.5σ , since 10.7'' at $r > 300''$ is at most 3.6 per cent in radius on a shallow gradient; and in the inner HST bins, where matched on bin edges the shifts reach 2.6σ and are mostly below 1.5. The underlying centre carries about 1'' of its own uncertainty, which is not propagated.

4 The dispersion estimator

The current HST, Gaia and Pristine measurements use the same mixture estimator, applied per annulus. Published MUSE dispersions are ingested separately. Diagnostic examples below record the data-preparation sequence; the current Gaia table and the final likelihood table identify the products frozen for rung 0.

Model. Each star contributes a two-component mixture. The cluster is a two-dimensional Gaussian whose principal axes follow that star’s own radial direction,

$$\mathbf{C}_i = R(\phi_i) \text{diag}(\sigma_R^2, \sigma_T^2) R(\phi_i)^\top + \mathbf{E}_i + \mathbf{D}_i, \quad (1)$$

with $\mathbf{E}_i$ the star’s proper-motion error covariance and $\mathbf{D}_i$ a rank-one term for the apparent dispersion induced by the cluster’s line-of-sight depth. For *Gaia* and Pristine, $\mathbf{E}_i$ carries the catalogue’s correlation coefficient; for HST it is diagonal, because the oMEGACat table supplies no correlation column. Any real HST correlation therefore leaks into the inferred σ_R and σ_T through the omitted $\pm 2\rho\epsilon_\alpha\epsilon_\delta\sin\phi\cos\phi$ terms; the size of that leak is not established here. The field is an *empirical* two-dimensional proper-motion density measured outside the cluster, not a Gaussian. The free parameters per annulus are σ_R , σ_T , the two mean motions, and the field fraction.

Why two dimensions. Projecting the field onto each star’s radial direction mixes its two unequal widths (3.1 and 2.0 mas yr^{−1} for ω Cen) around the annulus and manufactures shoulders that no one-dimensional model reproduces. Fitting in two dimensions removes the problem: the per-annulus χ^2 improved from 8.22 to 1.73 per bin when the field model was changed from two Gaussians to the empirical template.

Two defects, found and fixed. An external review of the implementation on 2026-09-19 found both. The covariance cross-terms in the mean-update normal equations carried a plus where the derivation gives a minus; they vanish for an isotropic cluster with uncorrelated errors, which every test mock was, so nothing caught it. On a synthetic anisotropic sample with an injected tangential mean of −0.200 mas yr^{−1} the code returned −0.125 and the corrected algebra −0.205, matching a direct numerical maximisation. The combined dispersion moved by at most 0.17 per cent, but the

anisotropy moved and the fitted rotation by up to 17 per cent. The independent check: our fitted rotation against the published curve the estimator never sees went from 0.93 of it to 0.98.

Separately, the profile-likelihood interval stepped by a fixed 0.02σ and returned the first point past the crossing, flooring *every* reported uncertainty at 1.4 per cent regardless of sample size; every HST bin sat exactly on it. It is now bracketed and interpolated, and recovers the analytic expectation to within 4 per cent from 2×10^3 to 10^5 stars. That fix is what exposed the rotation-uncertainty problem below.

Averaging over the stars, not the annulus. The model is averaged over eight equal-count quantiles of the radii of the stars actually measured in each bin, not over a complete annulus with $\Sigma(R)R$ weighting. The two differ where the selection’s coverage changes inside a bin: HST’s 300–340'' sample has a star-weighted mean radius of 313.2'' against the annulus midpoint of 320'', worth up to 1.17σ in the predicted dispersion.

Error deconvolution. Adding E_i to the intrinsic covariance *is* the convolution; the fit is unbinned in velocity and predicts the observed distribution directly. No subtraction is performed.

Systemic motion. The cluster’s three-dimensional velocity is projected onto each star’s own tangent basis following Vasiliev et al. (2020, §5.1), which removes the perspective contraction and the rotation of the equatorial basis across the field exactly. Removing nothing at all would leave an azimuthal pattern of amplitude $|\mu_{\text{sys}}| = 7.5 \text{ mas yr}^{-1}$, reading as a spurious dispersion of 5.3 mas yr^{-1} . Subtracting the constant equatorial vector removes that leading term; what the exact projection additionally removes is the residual, of order 0.06 mas yr^{-1} at 2400''.

Rotation. The fitted mean radial and tangential motions absorb streaming, so σ is measured about a rotating mean. The mean-square streaming is retained so that the Jeans model, which predicts the full second moment, is compared like with like. Our fitted mean tangential motion tracks the published rotation curve of Vasiliev & Baumgardt (2021) closely: 0.226 against 0.250 mas yr^{-1} at 435'', 0.159 against 0.177 at 773'', and 0.037 against 0.036 at 1722''.

5 HST

5.1 The catalogue and its quality flag

The oMEGACat VI catalogue contains 1,482,835 sources, 1,385,010 with proper motions, reaching 466''. Its published dispersion profile stops at 300''; the catalogue does not. The flag `selection_hq_astrometry` keeps 48 per cent of the stars with proper motions and requires, simultaneously: a temporal baseline longer than 10 yr; more than 75 per cent of individual measurements surviving clipping; reduced $\chi^2 < 5$ in both components; a proper-motion error inside the lower 95 per cent of the error distribution in its own 0.5-mag bin; and reliable photometry in *both* F625W and F814W. Stars fainter than $m_{\text{F625W}} = 24$ are cut outright, because there the error limit reaches 0.3 mas yr^{-1} , comparable to half the outer dispersion.

The two-filter photometric requirement is not cosmetic: F625W and F814W span 2002 to 2022, so demanding both certifies that the astrometry is supported across the whole baseline.

5.2 Our measurement and its validation

Applying the estimator of §4 to flagged stars only reproduces the published profile across the full radial range:

r (")	N	ours (mas yr ⁻¹)	published	ratio
3.8	146	0.7708 ± 0.0382	0.8048	0.958
8.3	595	0.7944 ± 0.0225	0.7854	1.012
18.0	3224	0.7654 ± 0.0108	0.7776	0.984
39.4	15574	0.7493 ± 0.0106	0.7440	1.007
85.9	66878	0.7070 ± 0.0100	0.7079	0.999
175.8	142615	0.6268 ± 0.0008	0.6236	1.005
270.2	119576	0.5511 ± 0.0008	0.5504	1.001
310.8	16912	0.5257 ± 0.0020	0.5199	1.011

Table 2: Median ratio 1.0042, scatter 2.5 per cent, the scatter arising in the innermost bins where only a few hundred stars pass the flag.

5.3 The outer limit

The flag survives to 360", on 66 stars in 340–360" and none beyond. The reason is specific: beyond 360" no star in the catalogue has F625W or F814W photometry at all, so the two-filter requirement rejects them automatically. This is not a threshold that can be relaxed.

Stars failing the flag are not usable as a substitute. Running the identical estimator on each subset where both exist gives an unflagged-to-flagged dispersion ratio of 1.092, 1.088, 1.079 and 1.052 at 150–200, 200–250, 250–300 and 300–340", a weighted mean of 1.077 ± 0.011 . This is a ratio of *dispersions*, not of errors: since intrinsic and measurement variances add, an excess of this size would require substantially larger error underestimates, and population differences and residual contamination would have to be excluded before attributing it to errors at all. It is used below only as an empirical sensitivity, never as a calibrated error correction in a fitted dataset. Taken raw, without a field model, they are worse still, and the apparent dispersion *rises* outwards (0.80, 0.88, 1.78 mas yr⁻¹ at 340–380, 380–420, 420–466"), because HST proper motions are relative to the cluster so field stars sit 7.5 mas yr⁻¹ away and dominate a raw variance even at a 1–2 per cent contamination level.

6 Gaia EDR3

6.1 Why the published profile is not used

Vasiliev & Baumgardt (2021) publish a proper-motion dispersion profile on a uniform 24" grid from 0 to 2400", obtained as a cubic spline in radius with 2–5 nodes, fitted jointly with a Plummer cluster profile and a two-Gaussian field by MCMC. Five of the eight points our likelihood originally took from it lie inside 380", at 24, 48, 96, 168 and 336". Their own quality flag passes *no* star inside 200" and five between 200 and 300". The inner points are therefore an inward continuation of a stiff global model, not a measurement. Its quoted uncertainties, as small as 0.002 mas yr⁻¹, are the posterior width of a 2–5 parameter spline constrained by 10⁵ stars, not the precision of any annulus.

6.2 The error model

The released catalogue carries the *raw* Gaia uncertainties. We verified this directly by joining all 228,055 members to `gaia_edr3.gaia_source`: the median ratio of catalogue to raw error is 1.0000, with 1st and 99th percentiles 0.9992 and 1.0009. The density-dependent inflation $\eta = (1 + \Sigma/\Sigma_0)^\zeta$ with $\zeta = 0.04$, $\Sigma_0 = 10$ (5-parameter) or 5 (6-parameter) is applied inside their fit, not in the file.

Deconvolving the raw errors therefore removes too little noise, and our first outer measurement ran 4–6 per cent high between 460 and 1000". Applying η removes the offset entirely. But η is not correct either: the authors' own non-circular test is that a correct error model makes the deconvolved dispersion independent of magnitude, and it is not. With raw errors the faint-to-bright

dispersion ratio is 1.114, 1.135 and 1.077 at 460–700, 700–1000 and 1000–1500''; with η it becomes 0.954, 0.980 and 0.951. Raw errors under-correct, the density-only η over-corrects.

6.3 Our measurement

We reduce sensitivity to this choice by selecting stars with small errors. The profile uses quality-flagged stars with $\frac{1}{2}(\epsilon_\alpha + \epsilon_\delta) < 0.4 \sigma_{\text{pub}}(R)$, where σ_{pub} is the published profile used only as a pilot. The intended bound is $\frac{1}{2}(0.4)^2(\eta^2 - 1) \simeq 2$ per cent; it is approximate rather than strict, because the threshold is applied to the arithmetic mean of the two component errors and against a pilot rather than the inferred dispersion, so an individual component may exceed it. The default measurement uses raw errors. The η fit is stored as a separate variant; half their separation is retained as a sensitivity diagnostic and is *not* added to the current statistical error. Neither the two endpoints nor their midpoint is established to bracket the true dispersion.

r (")	N	σ (mas yr ⁻¹)	stat	sys	σ_T/σ_R	f_{field}
357	50	0.4925	0.0352	0.0002	0.931	0.04
436	267	0.4440	0.0142	0.0007	0.839	0.06
532	1474	0.4365	0.0062	0.0017	0.859	0.06
662	2764	0.4036	0.0042	0.0019	0.870	0.10
830	3182	0.3627	0.0037	0.0016	0.940	0.17
1044	2891	0.3222	0.0040	0.0012	0.988	0.34
1328	2556	0.2917	0.0049	0.0008	1.081	0.59
1692	2612	0.2546	0.0071	0.0004	1.103	0.85
2149	3294	0.2314	0.0115	0.0003	1.176	0.96

Table 3: Current raw-error Gaia product, built 2026-09-20. The combined dispersion is shown for compactness; the likelihood uses the two components. “sys” is half the raw/eta separation, a diagnostic excluded from the default error. The ratio is projected and measured about fitted means; intrinsic anisotropy requires dynamical modelling.

The inner edge is 300'' because the flag passes five stars in total inside it. The current selected first two bins contain 50 and 267 stars. Earlier quality-flag counts quoted for different annuli or before the final error cut are not these fitted sample sizes. Their small errors make the innermost measurements relatively insensitive to the raw/eta choice.

7 Cross-validation

7.1 HST against *Gaia* in their overlap

HST’s flag survives to 360'' and *Gaia*’s begins at 300'', so the overlap is 300–360''. Three routes, differing in how HST’s rejected stars are used, agree:

Route	HST	<i>Gaia</i>	<i>Gaia</i> /HST
300–340'', both flagged	16912 at 311''	17 at 318''	1.028 ± 0.125
300–380'', both flagged	16978 at 311''	52 at 356''	0.934 ± 0.065
300–460'', HST corrected	92568 at 327''	328 at 430''	0.967 ± 0.029

The three routes are not independent: they share the same HST stars and nest in radius, so they are three treatments of one comparison rather than three confirmations. The first route is the cleanest: the two samples land at 311'' and 318'', essentially the same place, so nothing is extrapolated and no correction is applied on either side. Within an annulus the two samples do not sit at the same radius in general, because HST’s coverage falls outwards while *Gaia*’s flag passes more stars outwards; where that matters both are moved to the annulus midpoint on the local logarithmic slope before the ratio is formed.

A star-by-star match of the two catalogues is possible but informative only about the *unflagged* population, since no quality-flagged *Gaia* star has an HST counterpart. On identical stars, unflagged *Gaia* carries 0.90, 0.79, 0.59 and 0.44 mas yr^{−1} of scatter beyond its quoted errors at 0–100, 100–200, 200–300 and 300–460″, falling outwards exactly as crowding should.

7.2 A second selection: Pristine

The Pristine periphery catalogue provides a different survey, a different membership definition and a different field model. It is *not* an independent astrometric check: its proper motions are *Gaia* measurements, and the two samples share stars directly, 33 of the 52 selected *Gaia* stars in the innermost annulus and 11–37 per cent further out. A shared astrometric or estimator defect would leave the ratio near unity while both profiles were wrong, and the quoted ratio uncertainty treats the two as independent, which they are not. Selecting on metallicity alone ([Fe/H] < −1.2; the field median is −0.26 against −1.49 for the cluster) and applying the same estimator with a field template built from the same catalogue beyond 3.6°, with *no* proper-motion cut, gives a ratio to our *Gaia* profile on nine identical annuli from 9 to 57 pc of **1.005 ± 0.012**, with every bin within 5 per cent. Read as a test of the selection and the estimator, not of the astrometry, this is informative: the membership criterion, the magnitude range and the field model all change and the answer does not.

8 Supporting datasets

MUSE. 29 line-of-sight dispersion points from 5.7 to 290″, taken as published. Line-of-sight and proper-motion dispersions are permitted to differ, since they probe different projections of an anisotropic velocity ellipsoid.

***Gaia* DR2 — excluded.** Four of its nine published points (Baumgardt et al., 2019) lie inside 380″, the range whose EDR3 equivalent we removed as inward extrapolation. DR2 has its own, brighter selection and we hold no star list for it, so we cannot assess those points as we assessed EDR3’s. It was dropped from the fitted dataset on 2026-09-20 rather than carried on trust; it remains loadable by key.

Periphery spectroscopy. 157 flagged members from 0.58 to 3.15°, with no proper-motion selection and therefore free of the circularity that afflicts proper-motion measurements drawn from a proper-motion-selected catalogue. They give 6.12 ± 0.48 km s^{−1} at 67 pc and 8.36 ± 1.07 at 83 pc.

9 Beyond the catalogue edge

The *Gaia* member catalogue is truncated at 0.670° (63 pc), which is the authors’ retrieval radius, not a boundary. Computed with AGAMA in the McMillan (2017) potential for $M = 3.55 \times 10^6 M_\odot$ and the orbit integrated from the measured systemic velocity, the Jacobi radius is 90 pc at pericentre and 185 pc at the present galactocentric radius, so the data stop at $0.70 r_J(\text{peri})$. These two numbers are currently hard-coded rather than recomputed by the pipeline, and they use the spherically averaged tidal derivative $\Omega^2 - d^2\Phi/dr^2$; a flattened potential and a different angular-frequency convention would shift them. They are used only to orient the discussion, never in the inference.

A background-subtracted star count in the Pristine catalogue, using eight control windows in proper-motion space, detects the cluster at 78.1 ± 10.2 stars deg^{−2} at 57 pc (7.7σ), 15.9 ± 4.0 at 76 pc (4.0σ) and 2.5 ± 1.1 at 100 pc (2.2σ), falling as r^{-6} and consistent with the field beyond 130 pc in a circular average. The two spectroscopic points that carry the kinematic plateau therefore sit on a real overdensity; the third, at 178 pc from five stars, does not. Beyond 1.6° the 23 surviving candidates are not isotropic (axial Rayleigh $p = 0.016$ at position angle 122° measured from north

through east, along-axis excess $0.34 \pm 0.16 \text{ deg}^{-2}$ against -0.02 across), which is what tidal tails look like. Two caveats: the Rayleigh statistic accounts for fitting the axis, but the subsequent wedge comparison is selected on the same data, and the angular uniformity of the survey selection after the metallicity and proper-motion cuts has not been established.

10 The dataset that enters the likelihood

Key	Kind	N	Radial span
<code>hst_pm_radial_ours</code>	PM radial	21	4–342'' (0.11–9.0 pc)
<code>hst_pm_tangential_ours</code>	PM tangential	21	4–342''
<code>muse_los_dispersion</code>	LOS	29	5.7–290''
<code>gaia_edr3_ours_radial</code>	PM radial	9	356–2148''
<code>gaia_edr3_ours_tangential</code>	PM tangential	9	356–2148''

Table 4: 89 points. HST is now our own measurement from flagged stars rather than the published profile, which reaches 360'' instead of 300'' and so genuinely overlaps *Gaia*; and *Gaia* enters as its two components rather than their combination, so the anisotropy constrains the model where a halo would act. Fewer points than the previous 127, but each far better measured: the HST uncertainties are 0.1–4.5 per cent against a uniform 1.4 per cent before.

11 Known limitations

1. **The field template still differs from the science selection in its quality cuts.** Its parallax cut was removed on 2026-09-20 and the template re-queried (231 269 stars, up from 192 054), which moved the fitted dispersions by at most 0.33σ ; per-annulus matching on error ceiling and magnitude was added the day before. What remains unmatched is the template’s RUWE, excess-noise and PM-error ceilings against the science sample’s quality flag. Those are the same family of cuts, so the residual mismatch is smaller than the two already removed, but it is not zero.
2. **HST’s frame is not restored by adding a constant.** The catalogue’s proper motions are locally corrected, so position-dependent cluster streaming has been removed from them. Adding the constant systemic vector to score them against an absolute field template therefore leaves them displaced where the removed streaming was largest, of order 0.25 mas yr^{-1} . The numerical effect on the membership responsibilities has not been measured.
3. **The error cut uses a pilot profile.** The 0.4σ threshold is evaluated against the published profile, which is the object being replaced. The selection is on errors only and so does not truncate velocities, but its unbiasedness requires the quality selection to be independent of kinematics at fixed radius. ω Cen shows partial energy equipartition (Häberle et al., 2025), so selecting brighter stars selects more massive ones, which have a lower dispersion.
4. **Per-bin errors are treated as independent.** Shared stars between routes, a shared field template and spatially correlated *Gaia* astrometric systematics are not propagated. The raw/eta choice moves bins coherently and is now evaluated as separate fits; its half-separation is no longer added in quadrature. Propagated rotation errors are still treated independently across bins. Marginalisation over shared calibration or rotation uncertainties is not implemented.

12 Rotation

The Jeans model predicts the full second moment, so each dataset is compared with $\sqrt{\sigma_{\text{model}}^2 - \langle \bar{v}^2 \rangle}$. The streaming term is *uncertain*, and since 2026-09-20 that uncertainty is carried for our HST and Gaia profiles: an error δv moves the dispersion the model must predict by $v \delta v / \sigma$, which is added in quadrature to the measurement error. The MUSE rotation–dispersion covariance is unavailable and this propagation is not applied to MUSE. It matters most where the statistical errors are smallest — the HST tangential uncertainties at 223 and 270'' grow by factors of 4.9 and 5.7 — and it is a floor rather than the full correction, because the rotation curve moves coherently between bins and is here treated as independent. Where that streaming term comes from differs by instrument, and the difference is not cosmetic: it removes between 1 and 17 per cent of the squared dispersion in the *Gaia* bins.

Gaia offers two choices: the default published rotation curve, or our fitted per-annulus means. HST cannot supply its own absolute rotation from these proper motions. Its proper motions are *locally corrected*, with zero motion defined as the mean motion of the cluster stars in each patch, so the cluster’s streaming has been spent on setting the reference frame. Measured directly, our mixture recovers a residual mean tangential motion of 0.002–0.008 mas yr^{−1} over 40–340'' where the true rotation is 0.054–0.209: 96 per cent of it is gone. HST’s rotation term therefore comes from the [Vasiliev & Baumgardt \(2021\)](#) curve, and our own cannot replace it, because we fit rotation only outside 356'' where HST ends.

The two Gaia prescriptions change the model’s streaming correction and propagated uncertainty; they leave the measured dispersions unchanged. The figures `master_datasets_rotation_ours.png` and `master_datasets_rotation_published.png` compare reconstructed second moments under the two choices. Only Gaia differs: HST carries the published curve in both. These figures are sensitivity diagnostics, not completed variant mass fits.

13 Choices left explicit

Two modelling choices are not averaged over and not folded into a systematic. Each is a switch, and the intention is to fit both and compare the mass models rather than to guess:

1. **The *Gaia* error model** (`-gaia-errors`). The default is `raw`: the catalogue’s own uncertainties, unmodified. The alternative, `eta`, applies the density-dependent inflation of [Vasiliev & Baumgardt \(2021\)](#). Both fits are stored per component in the product. Until 2026-09-20 the two were averaged and half their separation carried as a per-bin systematic, which both hid the choice and treated one coherent calibration decision as independent noise.
2. **The rotation curve** (`-gaia-rotation`). The default is `published`: the [Vasiliev & Baumgardt \(2021\)](#) curve, with its published percentiles propagated into the error. The alternative, `ours`, removes the mean motion our own mixture fits in each annulus. The measured dispersions stay fixed; the streaming variance and its propagated uncertainty change. This is a sensitivity test, not an independent rotation dataset.

The calibration choices and rotation uncertainty can act coherently across bins while the likelihood is diagonal. Separate variant fits test sensitivity; they do not marginalise over that uncertainty. Reports now replay the saved likelihood arrays for new runs; legacy runs reload current products with their recorded Gaia options. Both paths check the saved best-sample likelihood before plotting (see `docs/CODE_ANALYSIS_AUDIT.md`). The CLI provides no HST rotation variant.




Figure 1: Every dataset in the likelihood, converted to km s^{-1} at $D = 5.43$ kpc so that proper motions and line-of-sight velocities share an axis. One symbol and colour per dataset throughout; the radial and tangential components of the same stars share that symbol and are distinguished by the marker fill. Middle: the projected component ratio used to diagnose anisotropy. Bottom: radial coverage and point count.




Figure 2: Left: HST measured from its own flagged stars, which reaches $360''$, against *Gaia*, which starts at $300''$. The published oMEGACat profile, which stops at $300''$, is shown faintly. Right: the comparison by three routes differing in their use of HST's rejected stars.




Figure 3: The rebuilt *Gaia* EDR3 measurement against the published spline and against our earlier version, which deconvolved the raw catalogue errors and ran 4–7 per cent high. The lower panel is the ratio to the published spline. This is a historical preparation figure: its flagged-star counts and intermediate selections differ from the final fitted sample tabulated in this note.




Figure 4: Left: the profile extended outwards with Pristine through the same mixture, with no proper-motion selection, against the same catalogue read through a fixed window. Right: Pristine over *Gaia* on identical annuli, 0.995 ± 0.014 .




Figure 5: The profile from the centre to 400 pc. The dispersion stops falling near 50–100 pc, where the pericentric Jacobi radius lies.




Figure 6: Left: cluster-like stars on the sky, selected on proper motion and CaHK metallicity, with the spectroscopic members overplotted. Right: their background-subtracted surface density, truncating near the Jacobi radius.

References

- Baumgardt, H., Hilker, M., Sollima, A., Bellini, A. 2019, MNRAS, 482, 5138. <https://ui.adsabs.harvard.edu/abs/2019MNRAS.482.5138B>
- Häberle, M., Neumayer, N., Clontz, C., et al. 2025, ApJ, *oMEGACat VI*. <https://arxiv.org/abs/2503.04903>, <https://doi.org/10.3847/1538-4357/adbe67>
- Kuzma, P. B., Ishigaki, M. N. 2025, MNRAS, 537, 2752. <https://arxiv.org/abs/2502.01135>, data <https://doi.org/10.5281/zenodo.14791430>
- Kuzma, P. B., et al. 2026, MNRAS, 550, *Disruption of a Giant*. <https://arxiv.org/abs/2605.23474>
- Vasiliev, E., Baumgardt, H. 2021, MNRAS, 505, 5978. <https://arxiv.org/abs/2102.09568>
- Vasiliev, E., Belokurov, V., Erkal, D. 2020, MNRAS, 501, 2279. <https://arxiv.org/abs/2009.10726>